

Oswald Blassou

AI engineer · LLM agents in production · Full-stack TypeScript

Paris, France / US · Remote only · Full-time, contract or freelance

oswaldblassou@gmail.com · linkedin.com/in/oswaldblassou · github.com/oswaldb22

SUMMARY

I design and ship generative-AI features to production: LLM agents, RAG, document extraction, evals and observability, and the full-stack TypeScript application around them. EPITA engineer, ex-SAP, bilingual French/English. Fixed overlap window on US Eastern afternoons.

EXPERIENCE

Freelance AI engineer — Nov 2025 – present

- Late 2025, startup engagement: designed LLM agents, built a RAG pipeline, and audited an existing AI solution.
- Designed and built a B2B SaaS with an AI agent, in production with paying customers.
- Stack: Next.js, React, TypeScript, PostgreSQL, Vercel AI SDK, OpenRouter, Langfuse, Vercel.

SAP — Software Engineer — Feb 2021 – Nov 2025

- Owned the architecture of the LLM agent within a team of four, working with teams in Germany and India to ship LLM-based features; the result was a 40% improvement in customer process efficiency.
- Designed and built an output-reliability API for LLM responses that cut completion errors by 50%.
- Earlier on the platform: moved RPA automation to the cloud with a containerized Node.js engine (startup 2 min → 20 s); built an event-driven update system that cut software rollouts from weeks to under 5 minutes; halved CI/CD time (40 → 20 min).

Audion — Software Engineer Intern — Sep 2019 – Jan 2020

- Moved a ~115 QPS gateway service to TypeScript (runtime errors -35%); CI/CD on GitHub Actions and Cloud Build (deploys 1 h 45 → 20 min); Terraform migration to Google Cloud.

WHAT I CAN DO FOR YOU

- **Agent in production:** add an AI agent or assistant to your product, from prototype to production (model routing, tool use, observability, fallbacks).
- **Reliability audit:** make a shipped LLM feature dependable (eval set, error rate, cost per call, observability), with a written report and the fixes.
- **Document pipeline:** extract existing documents into structured data, then generate compliant outputs (PDF, DOCX).
- Fixed-scope four-week builds welcome.

SKILLS

- **AI:** LLM agents, RAG, document extraction, prompt engineering, evals, LLM observability, Vercel AI SDK, OpenRouter, Langfuse, LangChain
- **Full-stack:** TypeScript, Node.js, Next.js, React, PostgreSQL, REST APIs, microservices, Docker, Kubernetes, Terraform, CI/CD, GCP, AWS, Vercel

EDUCATION AND LANGUAGES

- **EPITA**, Paris — Engineering degree / MSc in Computer Science, 2021. Major: Information Systems & Software Engineering; minor: Cloud Architecture
 - French (native), English (fluent)
-

Contracts through a French company (SASU); W-8BEN-E for US clients.